

# Data Mining

---

CS57300  
Purdue University

Bruno Ribeiro

January 22, 2018

# Differences Between Classification & Prediction

---

- Classification
  - Observing feature  $x$  predict discrete-valued label  $y$
- Prediction (point estimate)
  - Observing feature  $x$  predict of real-value label  $y$
  - Forecasting: predictions + confidence intervals
- Ranking prediction (rank estimate)
  - Predict item ranking
    - E.g.: Google results, Netflix recommendations

# Predictive modeling

---

- Data representation:
  - Training set: Paired attribute vectors and labels  $\langle y(i), \mathbf{x}(i) \rangle$   
or  
 $n \times p$  tabular data with label ( $y$ ) and  $p-1$  attributes ( $\mathbf{x}$ )
- Task: estimate a predictive function  $f(\mathbf{x}; \theta) = y$ 
  - Assume that there is a function  $y = f(x)$  that **maps** data instances ( $\mathbf{x}$ ) to labels ( $y$ )
  - Construct a model that approximates the mapping
    - Classification: if  $y$  is categorical
    - Regression: if  $y$  is real-valued

# Modeling approaches

# Learning predictive models

---

- Choose a **data representation**
- Select a **knowledge representation** (a “model”)
  - Defines a **space** of possible models  $M = \{M_1, M_2, \dots, M_k\}$
- Use **search** to identify “best” model(s)
  - Search the space of models (i.e., with alternative structures and/or parameters)
  - Evaluate possible models with **scoring function** to determine the model which best fits the data

# Scoring functions

---

- Given a model  $M$  and dataset  $D$ , we would like to “score” model  $M$  with respect to  $D$ 
  - Goal is to rank the models in terms of their utility (for capturing  $D$ ) and choose the “best” model
  - Score function can be used to search over ***parameters*** and/or ***model structure***
- Score functions can be different for:
  - Models vs. patterns
  - Predictive vs. descriptive functions
  - Models with varying complexity (i.e., number parameters)

# Predictive scoring functions

---

- Assess the quality of predictions for a set of instances
  - Measures **difference** between the prediction  $M$  makes for an instance  $i$  and the true class label value of  $i$

$$S(M) = \sum_{i=1}^{N_{test}} d[\underbrace{f(x(i); M)}_{\text{Predicted class label for item } i}, \underbrace{y(i)}_{\text{True class label for item } i}]$$

Diagram illustrating the components of the predictive scoring function  $S(M)$ :

- Sum over examples**: Indicated by an orange arrow pointing to the summation symbol  $\sum$ .
- Distance between predicted and true**: Indicated by a green arrow pointing to the distance function  $d$ .
- Predicted class label for item  $i$** : Indicated by a blue arrow pointing to the expression  $f(x(i); M)$ .
- True class label for item  $i$** : Indicated by a red arrow pointing to the expression  $y(i)$ .

# Predictive scoring functions

---

- Common score functions:

- Zero-one loss 
$$S_{0/1}(M) = \frac{1}{N_{test}} \sum_{i=1}^{N_{test}} I[f(x(i); M), y(i)]$$

$$\text{where } I(a, b) = \begin{cases} 1 & a \neq b \\ 0 & \text{otherwise} \end{cases}$$

- Squared loss

$$S_{sq}(M) = \frac{1}{N_{test}} \sum_{i=1}^{N_{test}} [f(x(i); M) - y(i)]^2$$

Careful with definition of class labels!  
Why?

- More later...
- Do we minimize or maximize these functions?



# Scoring functions


---

- Guide search *inside* of algorithms
  - Select path in heuristic search
  - Decide when to stop
  - Identify whether to include model or pattern in output
- Evaluate results *outside* of algorithms
  - Measure the absolute quality of model or pattern
  - Compare the relative quality of different algorithms or outputs

Where's the search?

# Find Parameters that Minimize Model Score (Error)

Usually we maximize (- score)



# Searching over models/patterns

---

- Consider a **space** of possible models  
 $M = \{M_1, M_2, \dots, M_k\}$  with parameters  $\theta$
- Search could be over model structures or parameters, e.g.:
  - **Parameters:** In a linear regression model, find the regression coefficients ( $\beta$ ) that minimize squared loss on the training data
  - **Model structure:** In a decision trees, find the tree structure that maximizes accuracy on the training data

# Optimization over score functions

---

- **Smooth** functions:
  - If a function is *smooth*, it is differentiable and the derivatives are continuous, then we can use gradient-based optimization
    - If function is *convex*, we can solve the minimization problem in closed form:  $\nabla S(\theta)$  using **convex optimization**
    - If function is smooth but non-linear, we can use iterative search over the surface of  $S$  to find a local minimum (e.g., hill-climbing)
- **Non-smooth** functions:
  - If the function is *discrete*, then traditional optimization methods that rely on smoothness are not applicable. Instead we need to use **combinatorial optimization**

---

## Linear Methods for Classification

# Motivation 1/2

---

- Given  $x$  features of a car (length, width, mpg, maximum speed,...)
- Classify cars into categories based on  $x$

## small car rentals ›



compacts  
economy car rentals

## medium car & SUV rentals ›



Coupes  
Sedans  
intermediate  
SUV rentals

## large car & SUV rentals ›



standard SUVs  
premiums  
luxury car rentals

## fuel efficient & hybrid ›



Green car rentals

## high occupancy car rentals ›



12-passenger vans  
mini vans  
premium SUV rentals

## reservable models ›



Corvettes  
Infinitis  
BMW's & more

## Motivation 2/2

---

- A person arrives at the emergency room with a set of symptoms that could possibly be attributed to one of three medical conditions. Which of the three conditions does the individual have?
- An online banking service must be able to determine whether or not a transaction being performed on the site is fraudulent, on the basis of the user's IP address, past transaction history, and so forth.
- On the basis of DNA sequence data for a number of patients with and without a given disease, a biologist would like to figure out which DNA mutations are deleterious (disease-causing) and which are not.



# Linear Discriminant Function (Two Classes)

- Two classes
- $x$  is a D-dimensional real-valued vector (set of features)
- $y$  is the car class

$$y_c = \begin{cases} 1 & , \text{ if car } c \text{ is "small"} \\ -1 & , \text{ if car } c \text{ is "luxury"} \end{cases}$$

- Find linear discriminant weights  $\mathbf{w}$

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Score function is the **squared error**

$$\sum_{c \in \text{TestDataCars}} (y_c - y(x_c))^2$$

- **Search algorithm: least squares algorithm**




Figure: C. Bishop

---

## How to Deal with Multiple Classes?

## Naïve Approach: one vs. many Classification

---

- How to classify objects into multiple types?

$$y_c^{(1)} = \begin{cases} 1 & , \text{ if car } c \text{ is "small"} \\ -1 & , \text{ if car } c \text{ is "luxury"} \end{cases}$$


$$y_c^{(2)} = \begin{cases} 1 & , \text{ if car } c \text{ is "small"} \\ -1 & , \text{ if car } c \text{ is "medium"} \end{cases}$$

$$y_c^{(3)} = \begin{cases} 1 & , \text{ if car } c \text{ is "medium"} \\ -1 & , \text{ if car } c \text{ is "luxury"} \end{cases}$$

Might work OK in some scenarios... but not clear in this case

## Issue with using binary classifiers for K classes

---



---

## Using Least Squares for Multiple Classes (Solution)

## Encoding the K Classes

---

$$\mathbf{I}_K = \underbrace{\left( \begin{array}{ccccc} 1 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & & & \ddots & \\ 0 & 0 & 0 & \cdots & 1 \end{array} \right)}_{i_1, i_2, \dots, i_K} \Bigg\}^K$$

$i_3$

- We start by encoding the classes as a **one-hot** binary coding scheme
- Class 3 is encoded as  $i_3$ , the 3<sup>rd</sup> column of identity matrix  $I_K$

# Linear Function


---

- Assign item with features  $\mathbf{x}$  to class  $k$  if  $y_k(\mathbf{x}) > y_j(\mathbf{x}), \forall j \neq k$

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0}$$

- The decision boundary is a (D-1) dimensional hyperplane

$$(\mathbf{w}_k - \mathbf{w}_j)^T \mathbf{x} + (w_{k0} - w_{j0}) = 0$$



# Least Squares Solution in Matrix Form

---

$$\mathbf{y}(\mathbf{x}) = \widetilde{\mathbf{W}}^T \widetilde{\mathbf{x}}$$

where

$$\widetilde{\mathbf{W}} = \begin{bmatrix} w_{10} & \cdots & w_{k0} & \cdots & w_{K0} \\ \mathbf{w}_1^T & \cdots & \mathbf{w}_k^T & \cdots & \mathbf{w}_K^T \end{bmatrix}$$

and  $\widetilde{\mathbf{x}} = (1, \mathbf{x}^T)^T$

The least squares solution is

Pseudoinverse

$$\widetilde{\mathbf{W}} = (\widetilde{\mathbf{X}}^T \widetilde{\mathbf{X}})^{-1} \widetilde{\mathbf{X}}^T \mathbf{T} = \widetilde{\mathbf{X}}^\dagger \mathbf{T}$$

where

$$\widetilde{\mathbf{X}} = \begin{bmatrix} 1 & \cdots & 1 \\ \mathbf{x}_1^T & \cdots & \mathbf{x}_n^T \end{bmatrix}$$

and

$$\mathbf{T} = [i_{y_1}, \cdots, i_{y_n}]$$

One hot encoding of class  $y_1$




---

Working with non-linearly separable data...

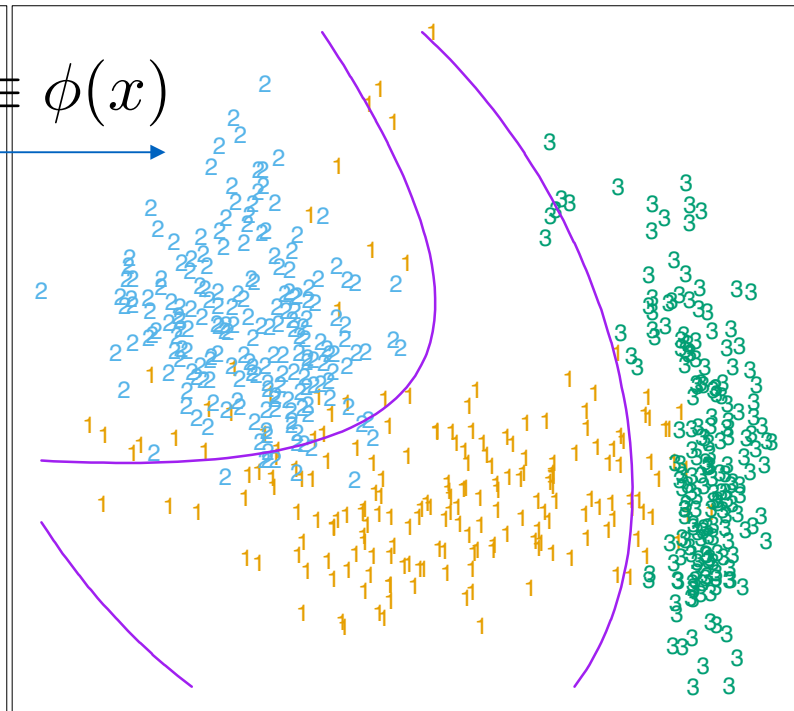
# Linear Methods → Linear Boundaries?

Data from three classes, with  
linear decision boundaries



Linear boundaries in the  
five-dimensional space  
 $\phi(\mathbf{x}) = (X_1, X_2, X_1X_2, X_1^2, X_2^2)$

$$\phi \equiv \phi(x)$$



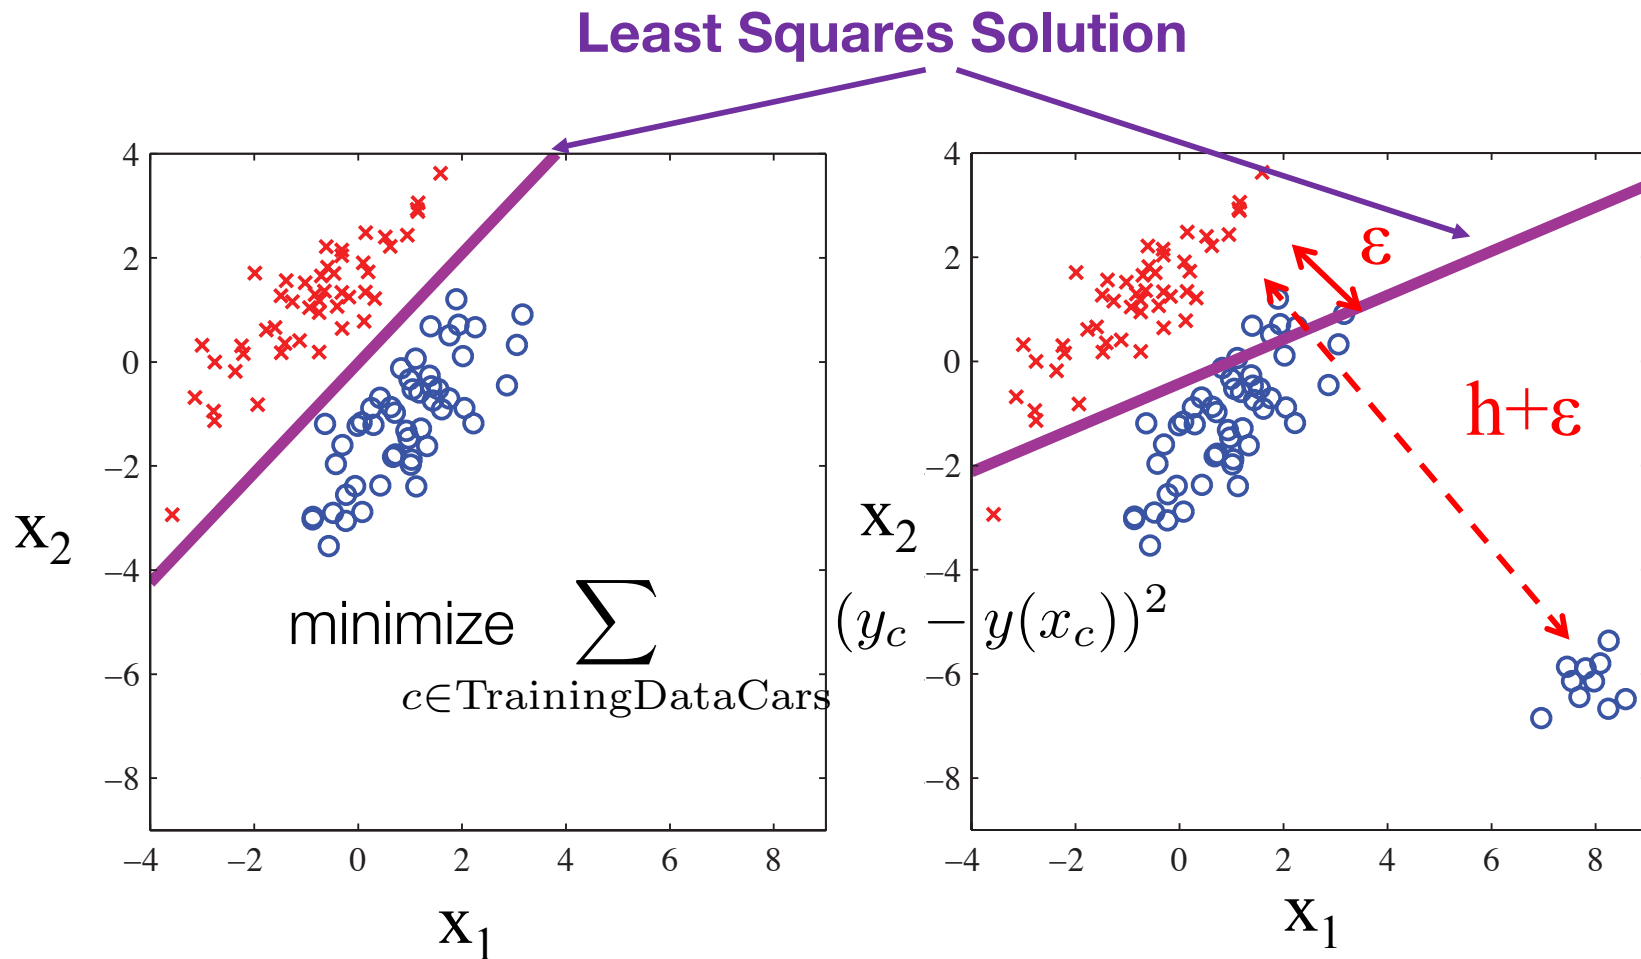
Linear inequalities in the transformed space are  
quadratic inequalities in the original space.

## Tips

---

- Helps if  $\phi_n = \phi(x_n)$  is **binary, quantized or normalized**
  - Example: if 1<sup>st</sup> feature  $x_n(1)$  is age of user  $n$ , then
    - $\phi_n(1)$  could be indicator if  $x_n(1)$  belongs to 1<sup>st</sup> quantile
    - $\phi_n(2)$  indicator whether  $x_n(1)$  belongs to 2<sup>nd</sup> quantile
    - ...
- Useful to add interaction terms  $\phi_h = \phi(x_n) \circ \phi(x_m)$ , where “o” is the Hadamard product (or element-wise product)
  - XOR operator can be better than Hadamard product for binary variables.

# Issues with Least Squares Classification



With square loss (score), optimization cares too much about reducing distance to boundary of well separable items...  
...solution is to change the score function


# Logistic Regression (for Classification)

---

- Back to two classes,  $C_1$  and  $C_2$
- Logistic regression is often used for two classes

$$p(C_1|\phi) = y(\phi) = \sigma(\mathbf{w}^T \phi)$$

where

$$\sigma(a) = \frac{\exp(a)}{1 + \exp(a)}$$


Logistic function

and  $\phi \equiv \phi(x)$



Transformation of feature vector (possibly non-linear)

- But can be easily generalized to  $K$  classes

## Finding Logistic Regression Parameters (binary classification)

---

- For a dataset  $\{\phi_c, t_c\}$  where  $\phi_c$  is transformed feature vector of car  $c$ ,  $t_c \in \{0,1\}$  is the class of car  $c \in \text{TrainingDataCars}$
- The likelihood function over the training data is

$$p(\mathbf{t}|\mathbf{w}) = \prod_{c \in \text{TrainingDataCars}} y(\phi_c)^{t_c} (1 - y(\phi_c))^{1-t_c}$$

where  $\mathbf{t} = (t_1, \dots, t_N)^T$ ,  $y(\phi_n) = \sigma(\mathbf{w}^T \phi_n)$

- The log of the likelihood function (log-likelihood) is

$$\log p(\mathbf{t}|\mathbf{w}) = \sum_{c \in \text{TrainingDataCars}} t_c \log y(\phi_c) + (1 - t_c) \log(1 - y(\phi_c))$$

- To solve the above equation, we **maximize the log-likelihood** over parameters  $\mathbf{w}$
- The above equation is also described as *cross-entropy loss*, *logistic loss*, *log loss*, on Tensorflow, Sklearn, and pyTorch.

# Solving Logistic Regression via Maximum Likelihood

---

- The above equations give the following gradient

$$\nabla_{\mathbf{w}} \log p(\mathbf{t}|\mathbf{w}) = \sum_{c \in \text{TrainingDataCars}} (t_c - y(\phi_c)) \phi_c = \mathbf{\Phi}^T (\mathbf{t} - \mathbf{y})$$

- The logistic function has derivative

$$\frac{d}{da} \sigma(a) = \sigma(a)(1 - \sigma(a))$$

Exercise: what are the dimensions of  $\mathbf{\Phi}$ ?

- The second derivative (Hessian) is then (verify!)

$$\mathbf{H} = \mathbf{\Phi}^T \mathbf{R} \mathbf{\Phi}$$

where  $R_c = y(\phi_c)(1 - y(\phi_c))$  and  $\mathbf{H}$  is a positive definite matrix. Because  $\mathbf{H}$  is positive definite the optimization is concave on  $\mathbf{w}$  and has a unique maximum.

## Iterative Update

---

- The iterative parameter update is (Newton-Raphson)

$$\begin{aligned}\mathbf{w}^{(\text{new})} &= \mathbf{w}^{(\text{old})} - (\Phi^T \mathbf{R} \Phi)^{-1} \Phi^T (\mathbf{y} - \mathbf{t}) \\ &= (\Phi^T \mathbf{R} \Phi)^{-1} \{ \Phi^T \mathbf{R} \Phi \mathbf{w}^{(\text{old})} - \Phi^T (\mathbf{y} - \mathbf{t}) \} \\ &= (\Phi^T \mathbf{R} \Phi)^{-1} \Phi^T \mathbf{R} \mathbf{z}\end{aligned}$$

where  $\mathbf{z}$  is an  $N$ -dimensional vector with elements

$$\mathbf{z} = \Phi \mathbf{w}^{(\text{old})} - \mathbf{R}^{-1} (\mathbf{y} - \mathbf{t})$$



# Training Logistic Regression with Data Selection Bias

---

- Let  $\pi_i$  be the probability of sampling example  $i$  in the training data
- We say a data sample is biased when  $\pi_i$  is **not** uniform
- Correcting for bias in the likelihood function:

$$\log p(\mathbf{t}|\mathbf{w}) = \sum_{c \in \text{TrainingDataCars}} \frac{1}{\pi_c} (t_c \log y(\phi_c) + (1 - t_c) \log(1 - y(\phi_c)))$$

where  $C_c$  is the class of car  $c$ .

- Generally, it is a good idea to test training data for different weights
  - The weights can be used to protect against sampling designs which could cause selection bias.
  - The weights can be used to protect against misspecification of the model.
- Unbalanced datasets:
  - Suppose there are more males than females in dataset. What will happen to decision boundary? How to fix it?

# Multiclass Logistic Regression (MLR)

---

- Consider  $K$  classes and  $N$  observations
- Let  $C_i$  be the class of the  $i$ -th example with feature vector  $\phi_i$

$$P(C_c = t_c | \phi_c) = \frac{\exp(\mathbf{w}_k^T \phi_c)}{\sum_{h=1}^K \exp(\mathbf{w}_h^T \phi_c)} \quad \text{a.k.a. } \mathbf{softmax}$$

- If we assume one-hot encoding of target variable  $t_c$ , the log-likelihood function is

$$\begin{aligned} & \sum_{c \in \text{TrainingDataCars}} \sum_{k=1}^K t_{c,k} \log \frac{\exp(\mathbf{w}_k^T \phi_c)}{\sum_{h=1}^K \exp(\mathbf{w}_h^T \phi_c)} \\ &= \sum_{c \in \text{TrainingDataCars}} \sum_{k=1}^K t_{c,k} \mathbf{w}_k^T \phi_c - \sum_{c \in \text{TrainingDataCars}} \log \sum_{h=1}^K \exp(\mathbf{w}_h^T \phi_c) \end{aligned}$$