

Data Mining & Machine Learning

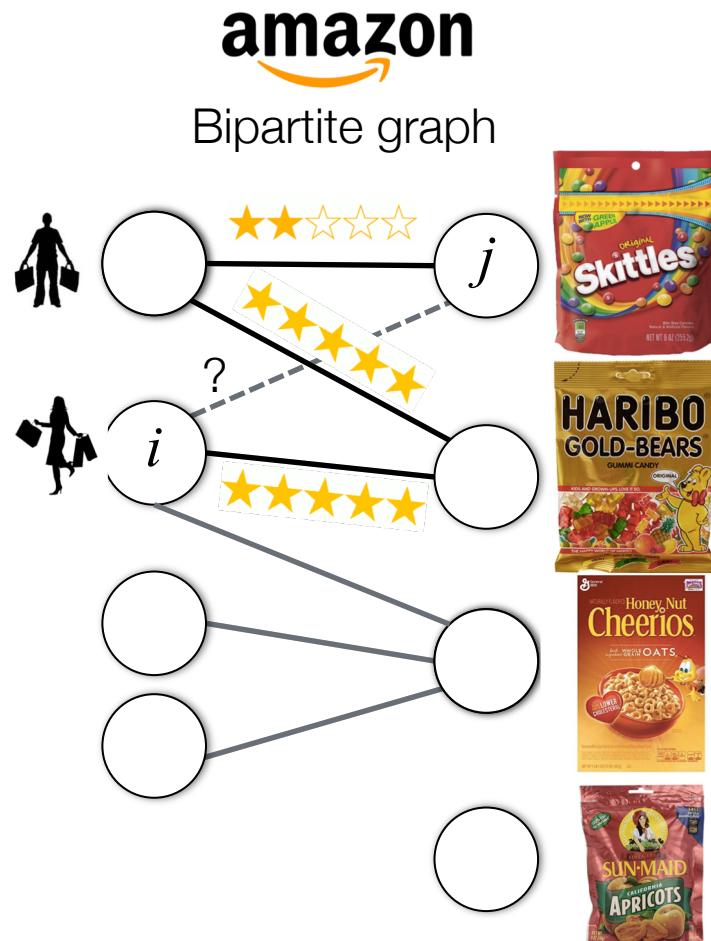
CS57300
Purdue University

April 3, 2018

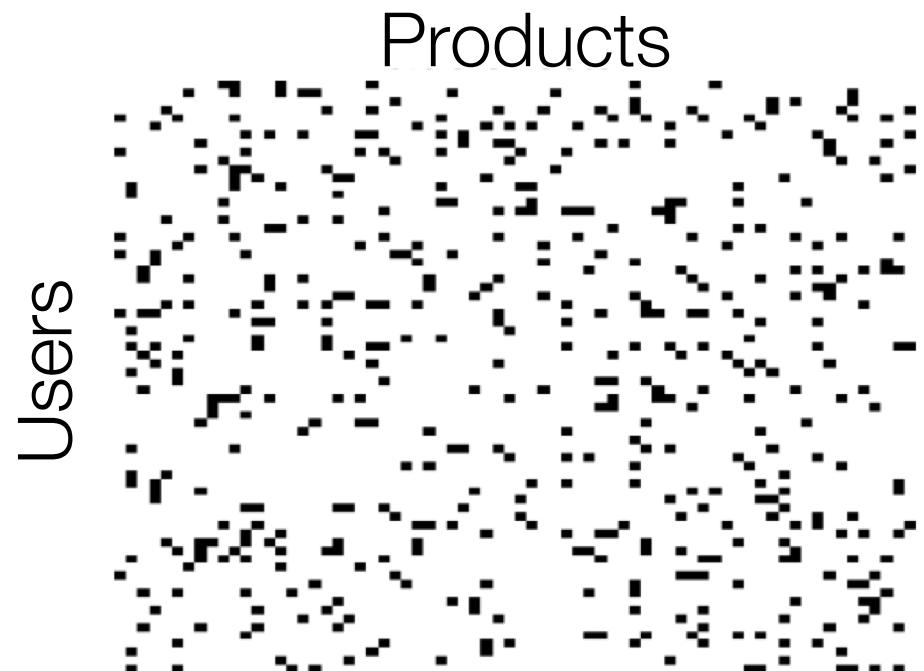
Collaborative Filtering

Product Recommendation

Example of collaborative filtering application.
Product recommendation based on user ratings



X is an $n \times m$ matrix.



$X_{i,j}$ - Rating user i gives to product j

Rating Matrix

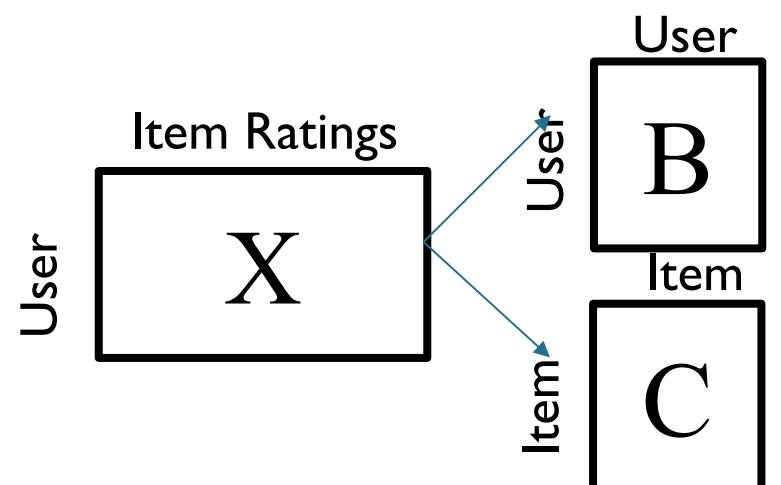
X is an $n \times m$ matrix.

	b_1	b_2	b_3	b_4	b_5	b_6
 a_1	x_{11}	x_{12}	?	?	x_{15}	x_{16}
a_2	?	?	?	x_{24}	?	?
a_3	?	x_{32}	?	?	?	?
 a_4	?	?	x_{43}	?	?	x_{46}
a_5	?	?	?	?	x_{55}	?

- ▶ Description of rating matrix should be invariant to permutation
 - Order of users (a) and products (b) in the matrix carries no information
- ▶ Absence / presence of rating is a signal
 - But means two things:
 - (a) user has never seen the product
 - (b) user does not want to rate the product
- ▶ Can be extended to graphs with vertex attributes

Simplest Idea: Recommendation via Similarity

- ▶ (user,user) similarity to recommend items
 - $B = XX^T / \langle \text{normalization factor} \rangle$
 - good if item base is smaller than user base
 - good if item base changes rapidly
- ▶ (item,item) similarity to recommend new items that were also liked by the same users
 - $C = X^T X / \langle \text{normalization factor} \rangle$
 - good if the user base is small
- ▶ Oldest known collaborative filtering method



Recommendations Can be Based on Similarity

$$\text{Heuristic: } X_{i,j} \approx \frac{\sum_{i,b} B_{i,b} X_{b,j}}{\sum_{i,b} B_{i,b}}$$

	items											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		2	5			5		4	
2			5	4			4			2	1	3
3	3	2	4		1	2		3		4	3	5
4		2	4		5			4			2	
5			4	3	4	2				2		5
6	6	1		3		3		2			4	



User-user similarity

$$B_{13} = 0.2$$

$$B_{16} = 0.3$$

weighted
average

$$\frac{0.2 \cdot 2 + 0.3 \cdot 3}{0.2 + 0.3} = 2.6$$

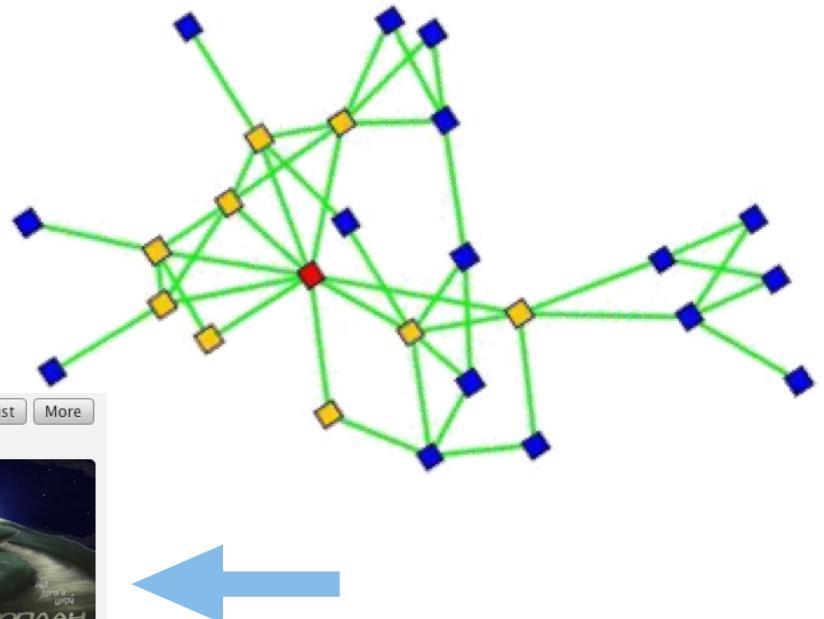
Properties

- ▶ Intuitive
- ▶ No (substantial) training
- ▶ Handles new users / items
- ▶ Easy to explain to user

Recommended for you

Image	Title	Artist	Description
	Casually Introducing	Walter Smith III	Similar to Eric Harland
	Companeros De Mi Vida	Eliades Ochoa	Similar to Cachao and Irakere
	Tibiri Tabara	Sierra Maestra	You've scrobbled Sierra Maestra, but not this release
	New York Ska-Jazz Ensemble	New York Ska-Jazz Ensemble	You've scrobbled New York Ska-Jazz Ensemble,
	More Late Night Transmissions With...	Jaya the Cat	You've scrobbled Jaya the Cat, but not this release
	Appetite For Destruction	Guns N' Roses	You've scrobbled Guns N' Roses, but not this release

+ Add as Playlist More

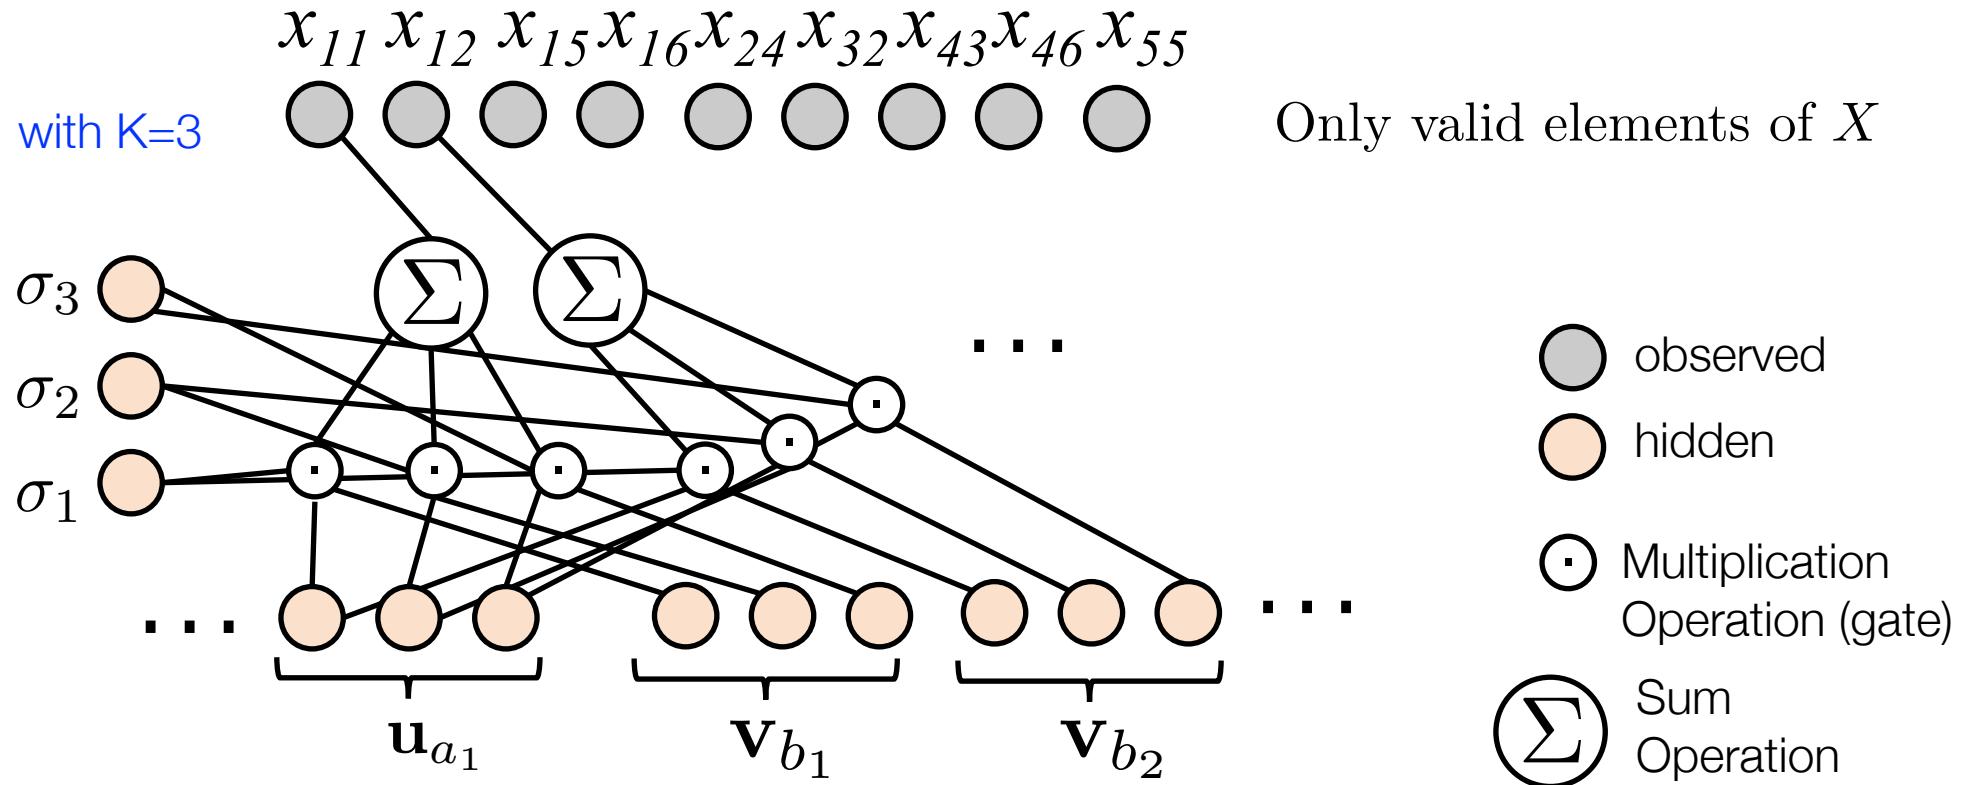


- ▶ Accuracy & scalability questionable

We will now propose something closer to the
state-of-the-art
(it is secretly a probabilistic model)

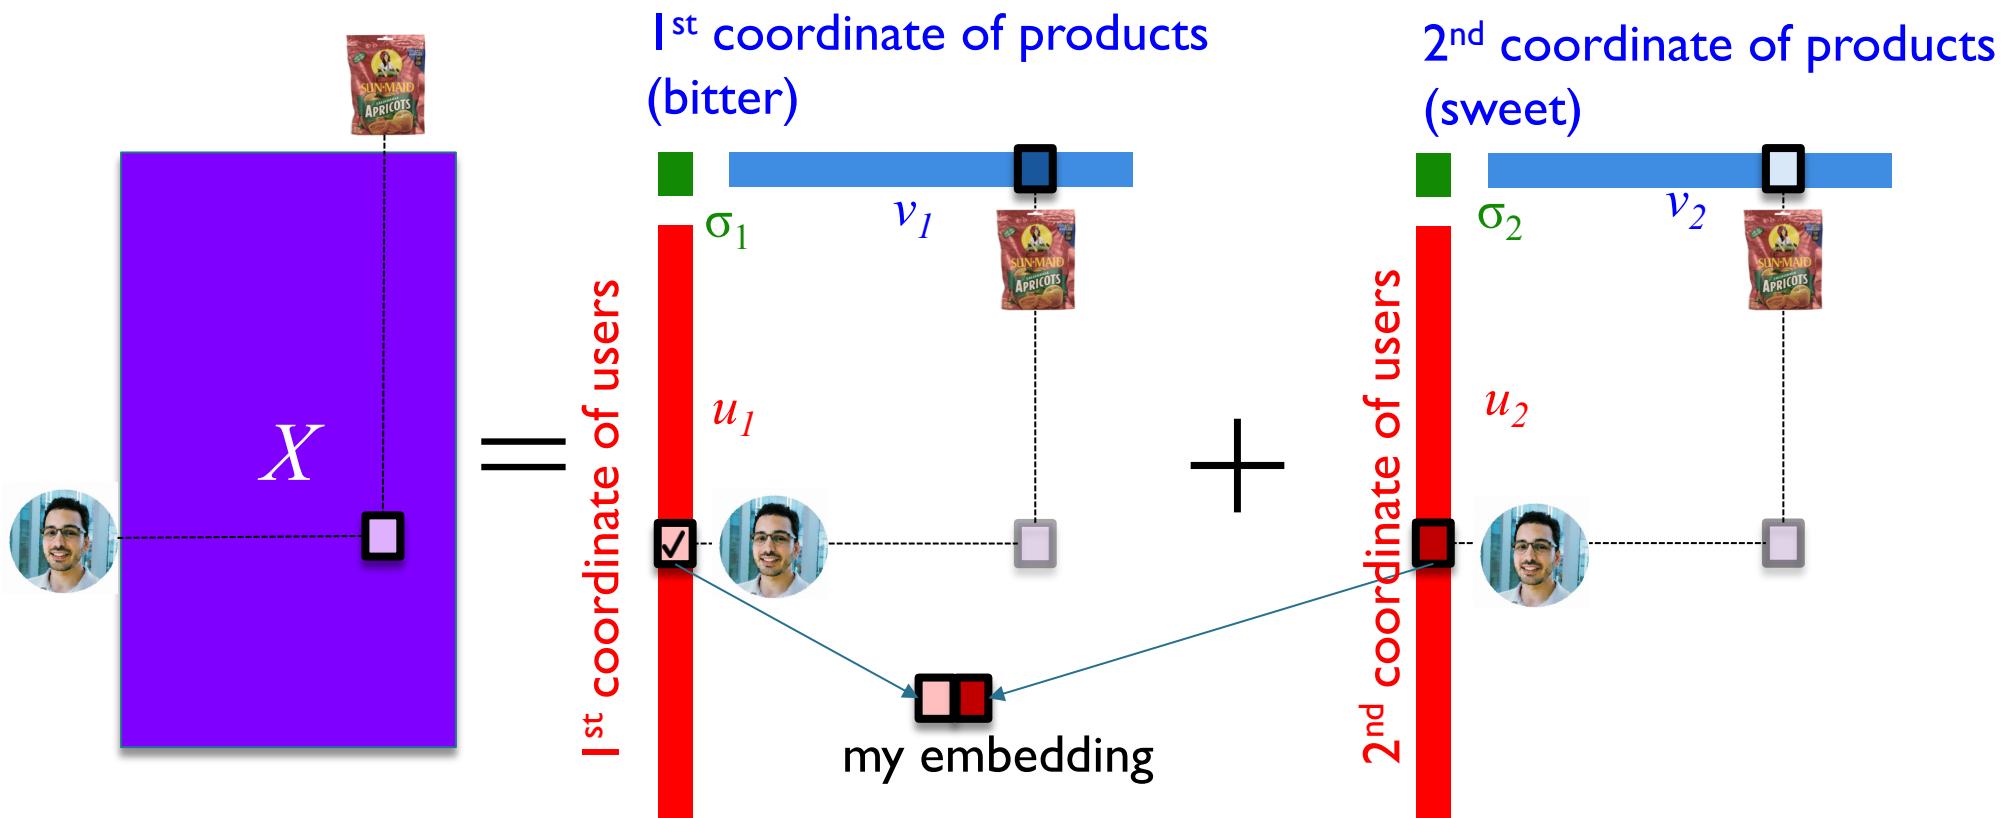
A Rating Matrix Model

- Assume X has is a set of random variables organized in a matrix (observable)
 - The model has the following parameters
 - \mathbf{v}_j is a $1 \times K$ vector of latent factors of product j
 - \mathbf{u}_i is a $1 \times K$ vector of latent factors of user i
 - σ_k is weight of dimension $k = 1, \dots, K$



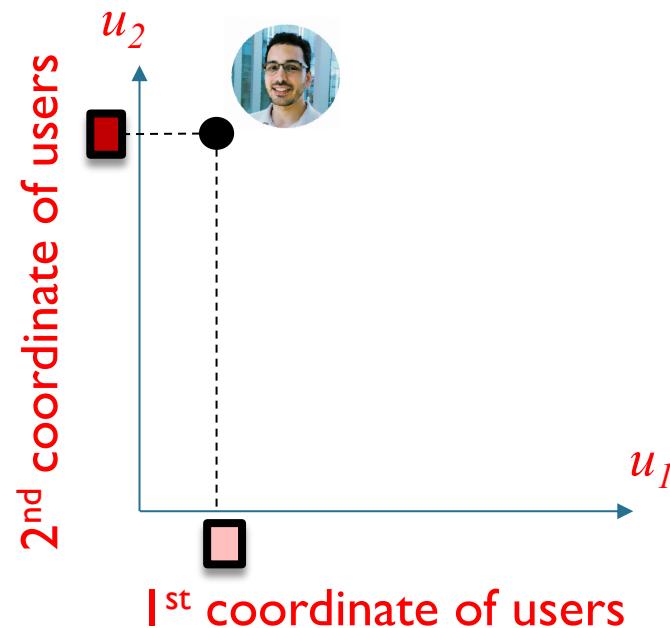
Alternative View: Matrix Decomposition

Illustration of a rank- K matrix approximation (assuming $\text{rank}(\mathbf{X}) < K$) :



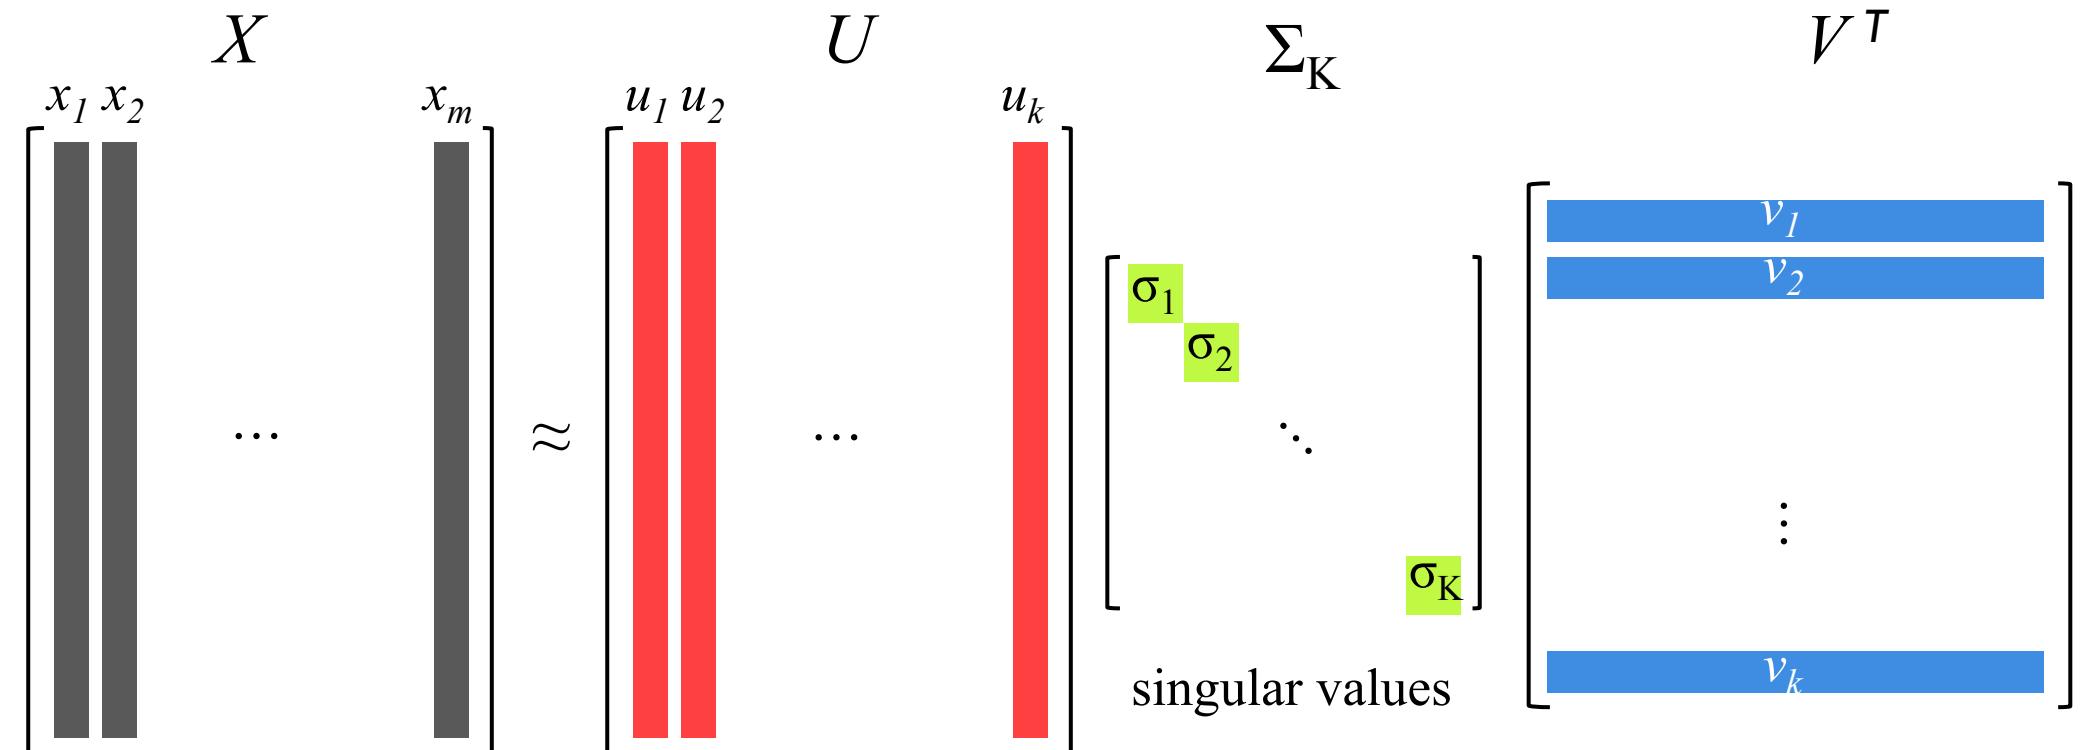
$$X \approx U\Sigma V^T = \sum_{k=1}^K \sigma_k u_k v_k^T$$

My Embedding (Representation as Latent Variables)



Singular Value Decomposition

$$\mathbf{X} \approx \mathbf{U}\Sigma_K\mathbf{V}^T$$



Data

Left singular vectors
(Columns are Orthogonal
and Normalized)

Right singular vectors
(Columns are Orthogonal
and Normalized)

Understanding SVD singular vectors

- ▶ Let $K = \text{rank}(X)$
- ▶ As U and V have *orthonormal** columns

**orthonormal* = orthogonal with norm 1

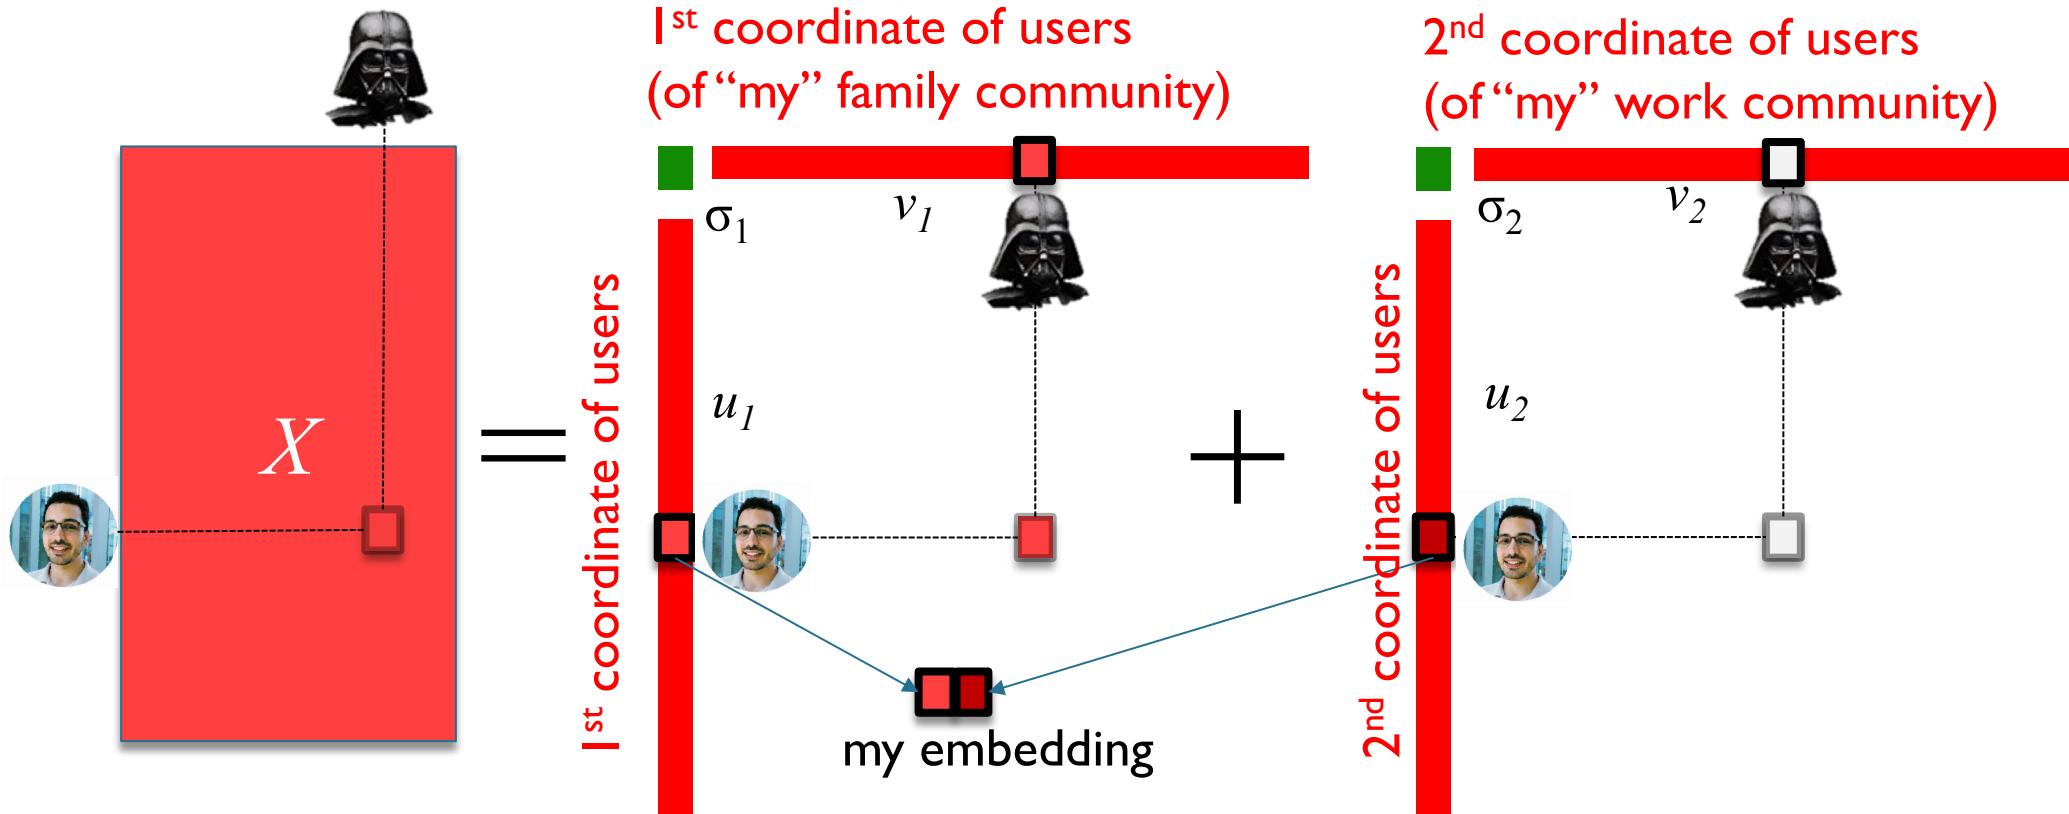
$$X^T X = (U \Sigma V^T)^T (U \Sigma V^T) = V \Sigma^2 V^T$$

$$XX^T = (U \Sigma V^T)(U \Sigma V^T)^T = U \Sigma^2 U^T$$

What do V and U represent?
(hint, slide 5)

Undirected Network Embedding

- Illustration of a rank-2 approximation:



$$X \approx U\Sigma U^T = \sum_{k=1}^K \sigma_k u_k u_k^T$$

iPython Notebook

```
from sklearn.decomposition import TruncatedSVD  
import matplotlib.pyplot as plt  
import numpy as np
```

```
maxi = 0  
maxj = 0  
Xlist = []  
# data at https://www.cs.purdue.edu/homes/rib/  
# Original data from http://snap.stanford.edu/datasets/  
with open("FBmatrix.csv") as f:
```

```
    for line in f.readlines():  
        i,j,Xij = list(map(int, line.split()))  
        maxi = max(i,maxi)  
        maxj = max(j,maxj)  
        Xlist.append((i,j,Xij))  
X = np.zeros((maxi+1, maxj+1))  
for i,j,Xij in Xlist:  
    X[i,j] = Xij
```

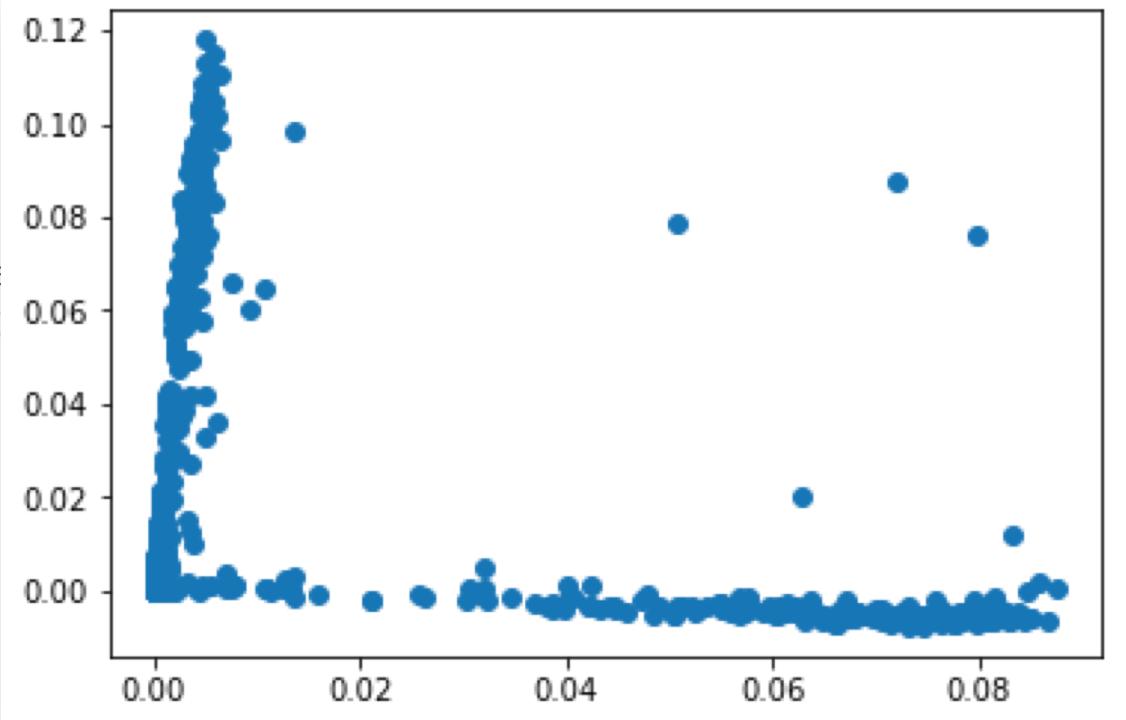
```
SVD = TruncatedSVD(n_components=2, n_iter=10)  
SVD.fit(X)  
# 2D embedding of FB egonet 1912
```

```
plt.plot(SVD.components_[0], SVD.components_[1], 'o')  
plt.show()
```

Creates
the adjacency
matrix

Finds latent
embedding

Plots the
embedding



2 dimensions

SVD: Optimization View

- ▶ Objective Function does not need Σ :

$$\operatorname{argmin}_{\mathbf{U}, \mathbf{V}} \sum_{i,j} (X_{i,j} - (\mathbf{u}_{i,*})(\mathbf{v}_{j,*})^T)^2 + \lambda(\|\mathbf{U}\|_2^2 + \|\mathbf{V}\|_2^2)$$

Row vectors of \mathbf{U} and \mathbf{V}

Regularization

- ▶ Can be easily implemented via alternating least squares

- ▶ What is the statistical assumption about \mathbf{X} ?

What does the L_2 minimization imply?

- Conditionally independent, Normally distributed ratings

$$X_{i,j} \sim \text{Normal}((\mathbf{u}_{i,*})(\mathbf{v}_{j,*})^T)^2, 1)$$

- But ratings are not Normal!
 - Ratings 1,...,5 or binary { , }

Poisson Matrix Factorization

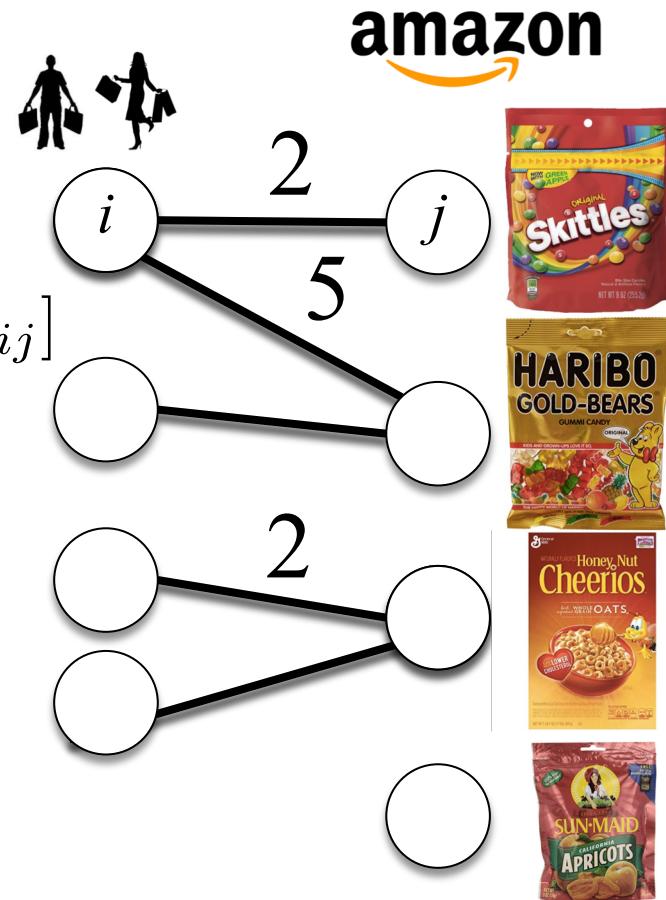
Users have a propensity rate to buy certain category of product (**W**)

In a category some products have a propensity rate to be bought (**H**)

$$\mathbf{X}_{ij} \sim \text{Poisson}([\mathbf{WH}]_{ij})$$

$$\text{MLE} \rightarrow \underset{\mathbf{W}, \mathbf{H}}{\operatorname{argmax}} \sum_i \sum_j [X_{ij} \log(\mathbf{WH})_{ij} - (\mathbf{WH})_{ij}]$$

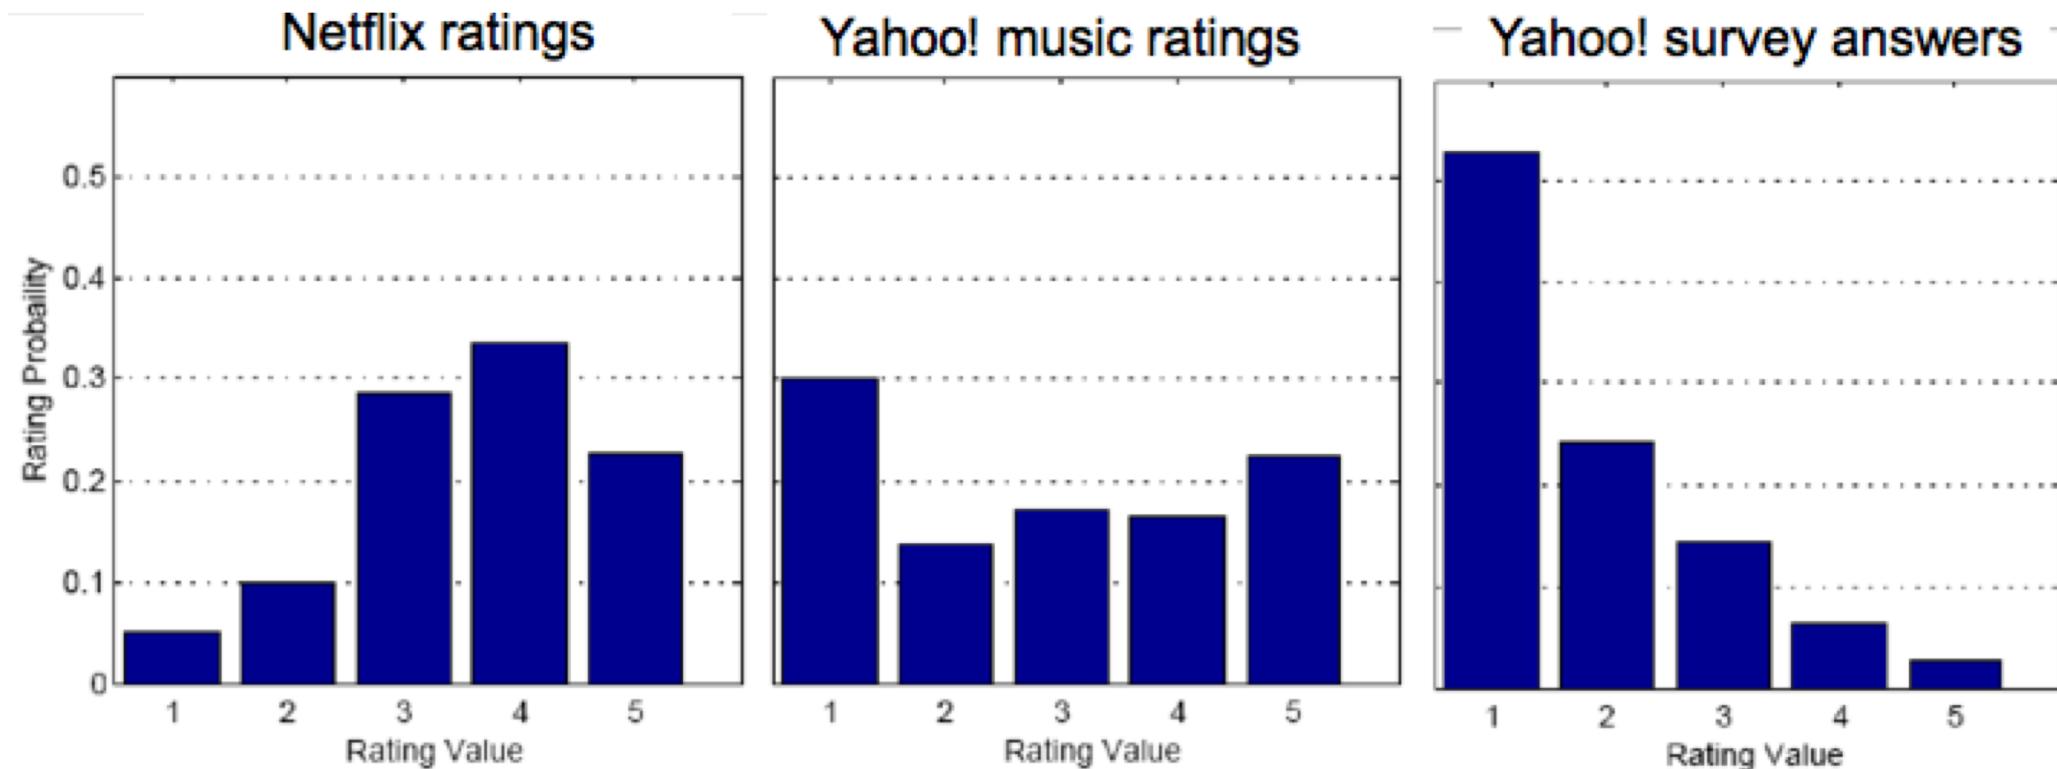
(Poisson Matrix Factorization,
Non-negative matrix factorization)



Data Quality Problems

- ▶ Amatriain et al. (2009): Asking users to **rate** items **again** reduce prediction error by as much as 14%
- ▶ Marlin et al. (2007): most existing algorithms based on the flawed assumption that missing ratings (data) are missing at random (uniformly)

Ratings are not given at random



- ▶ Marlin et al. "Collaborative Filtering and the Missing at Random Assumption" UAI 2007

Beyond Matrices

Adding Node/Edge Covariates is Generally Easy

- ▶ E.g.: RESCAL tensor factorization (a generalization of matrix factorization) (Nickel et al. 2011)

